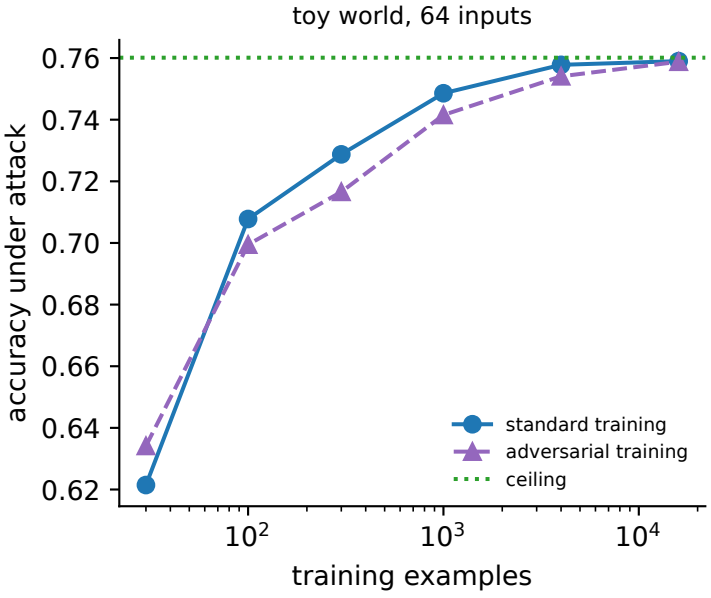
toy world, 64 inputs
accuracy under attack
0.76
0.74
0.72
0.70
0.68
0.66
0.64
0.62
10^2
10^3
10^4
training examples
standard training
adversarial training
ceiling

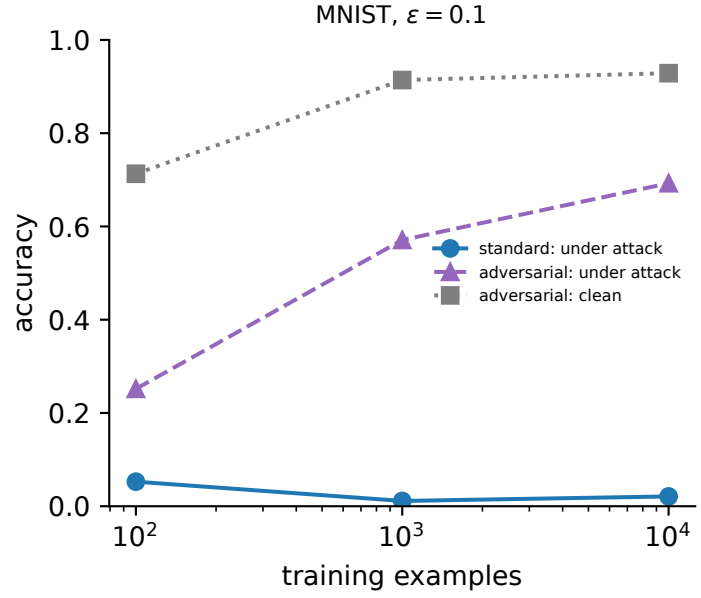
MNIST, $\epsilon = 0.1$
accuracy
1.0
0.8
0.6
0.4
0.2
0.0
10^2
10^3
10^4
training examples
standard: under attack
adversarial: under attack
adversarial: clean

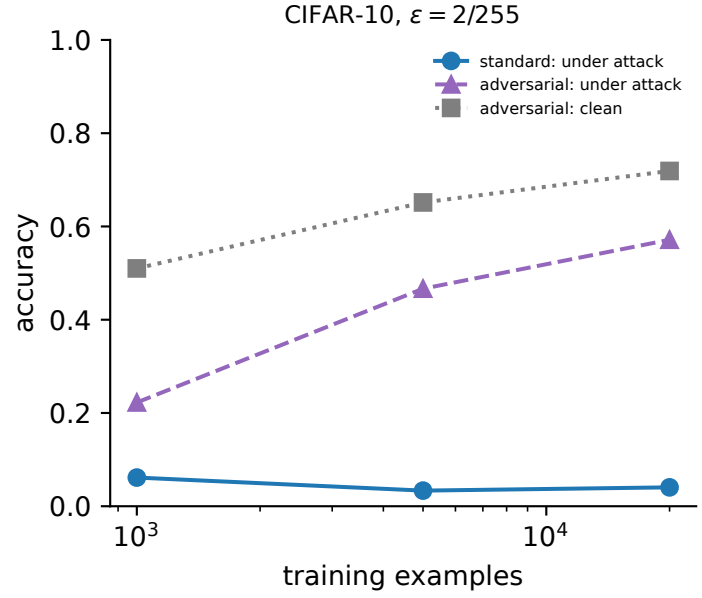
CIFAR-10, $\epsilon = 2/255$
accuracy
1.0
0.8
0.6
0.4
0.2
0.0
10^3
10^4
training examples
standard: under attack
adversarial: under attack
adversarial: clean