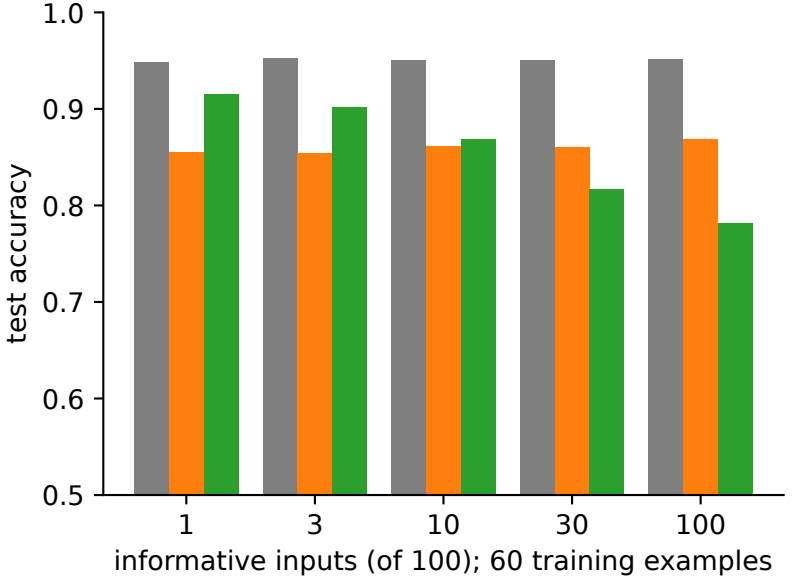
1.0
0.9
0.8
0.7
0.6
0.5
test accuracy
1
3
10
30
100
informative inputs (of 100); 60 training examples

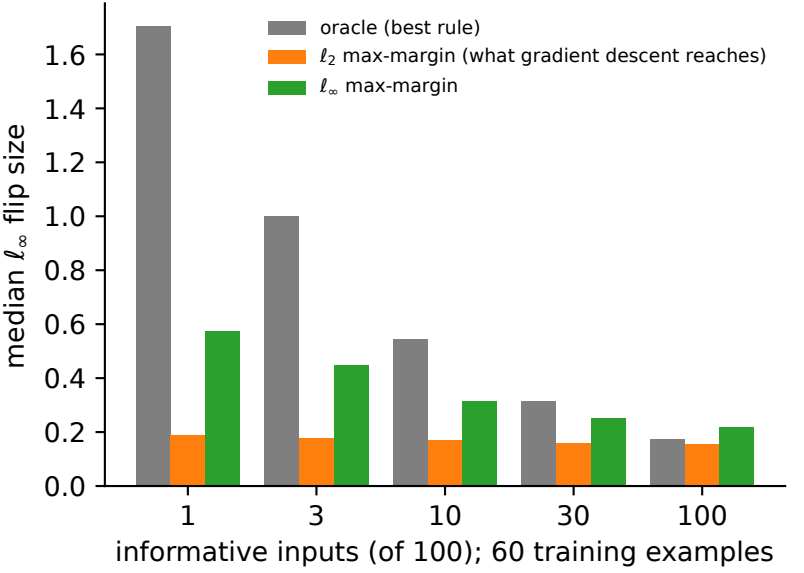
oracle (best rule)
ℓ_2 max-margin (what gradient descent reaches)
ℓ_∞ max-margin
1.6
1.4
1.2
1.0
0.8
0.6
0.4
0.2
0.0
median ℓ_∞ flip size
1
3
10
30
100
informative inputs (of 100); 60 training examples