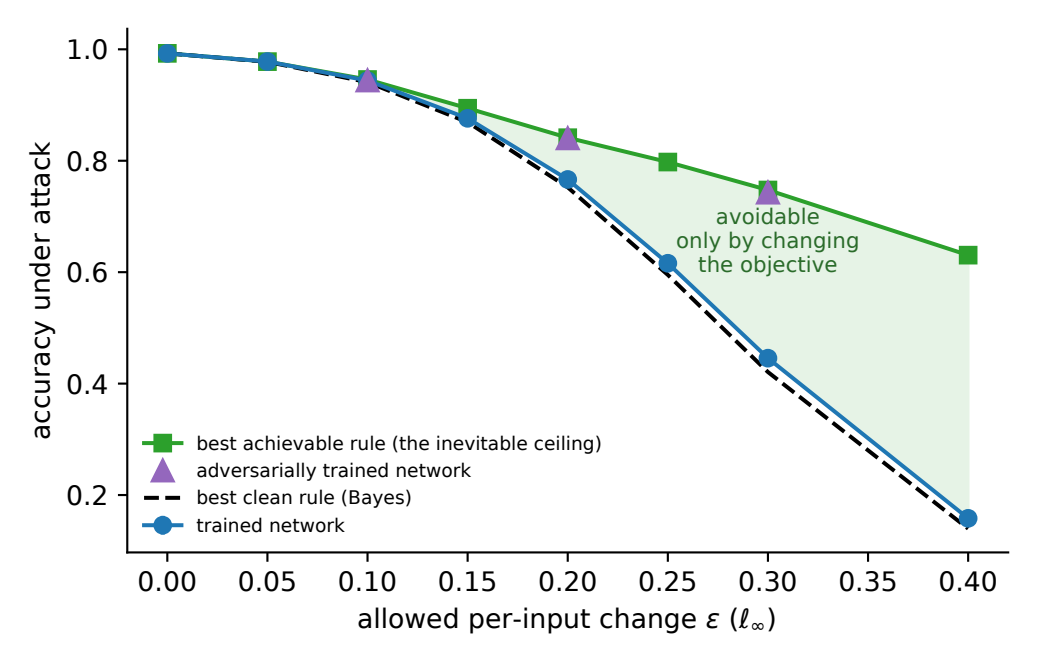
1.0
0.8
0.6
0.4
0.2
accuracy under attack
avoidable
only by changing
the objective
best achievable rule (the inevitable ceiling)
adversarially trained network
best clean rule (Bayes)
trained network
0.00
0.05
0.10
0.15
0.20
0.25
0.30
0.35
0.40
allowed per-input change ε (ℓ∞)