hypothesis	kind	toy world	real data	note
H1 evidence spread over many inputs	inevitable	supported	partly	geometry holds; MNIST already spread
H2 small, not unreliable; features carry signal	inevitable	supported	supported	transfer works on MNIST, weaker CIFAR-10
H4a irrelevant inputs ignored badly	avoidable	supported	supported	about 2x in toy, MNIST, wide MLP
H4b weight sparsity as a fix	avoidable	supported	not tested	
H4c activity sparsity	none	not supported	not tested	
H8 noise vs worst case	inevitable	supported	partly	right order of magnitude
H10 objective (adversarial training)	avoidable	supported	supported	closes the toy gap; costs accuracy
H7 optimiser bias (l2 vs l_inf margin)	avoidable	supported	supported	+50% flip size on MNIST, F-MNIST
H6 data quantity	none/avoidable	not supported	not supported	none (standard); strong (adversarial)
H5 off-manifold directions	avoidable	partly	partly	CIFAR-10 only, with a caveat
H3 shape vs texture / frequency	inevitable + arch.	partly	not supported	spread, not frequency, matters
H9 architecture / invariance	avoidable	supported	partly	helps MNIST, not consistently