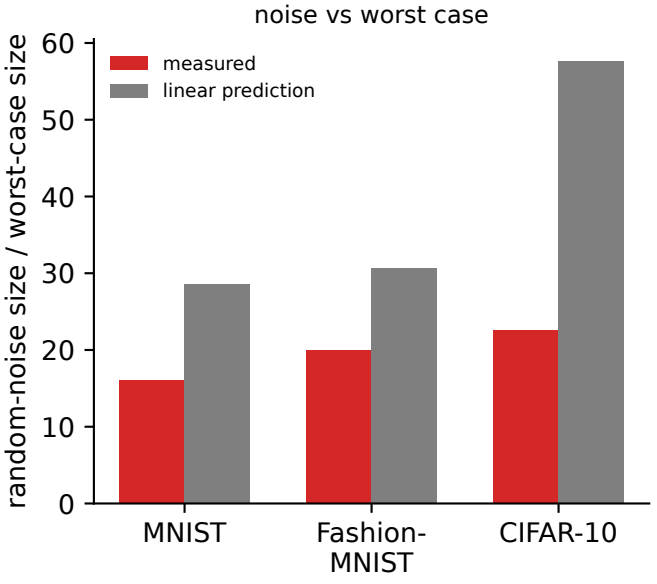
noise vs worst case
random-noise size / worst-case size
measured
linear prediction
0
10
20
30
40
50
60
MNIST
Fashion-MNIST
CIFAR-10

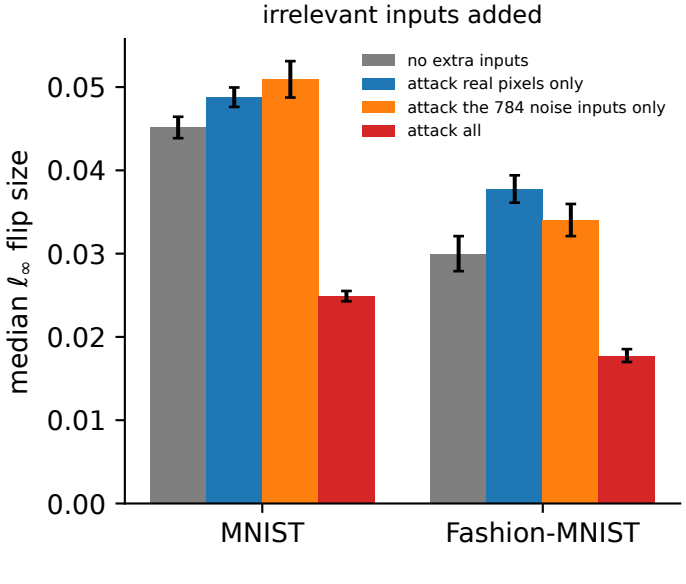
irrelevant inputs added
median ℓ_∞ flip size
no extra inputs
attack real pixels only
attack the 784 noise inputs only
attack all
0.00
0.01
0.02
0.03
0.04
0.05
MNIST
Fashion-MNIST

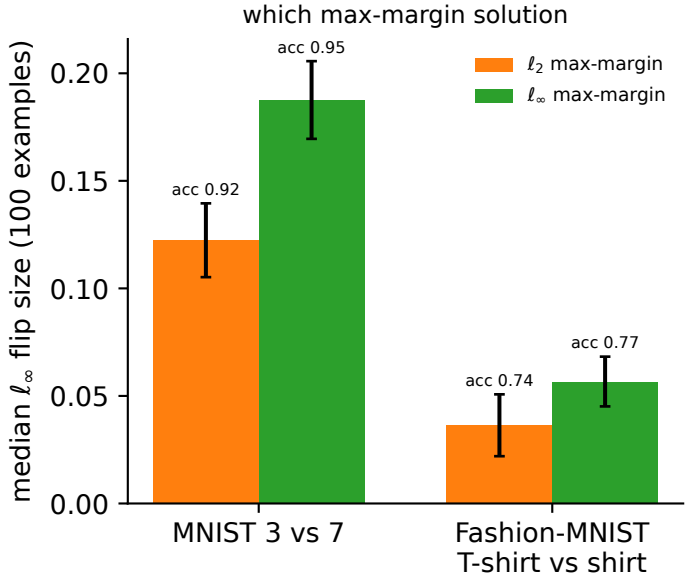
which max-margin solution
median ℓ_∞ flip size (100 examples)
ℓ_2 max-margin
ℓ_∞ max-margin
acc 0.92
acc 0.95
acc 0.74
acc 0.77
0.00
0.05
0.10
0.15
0.20
MNIST 3 vs 7
Fashion-MNIST
T-shirt vs shirt