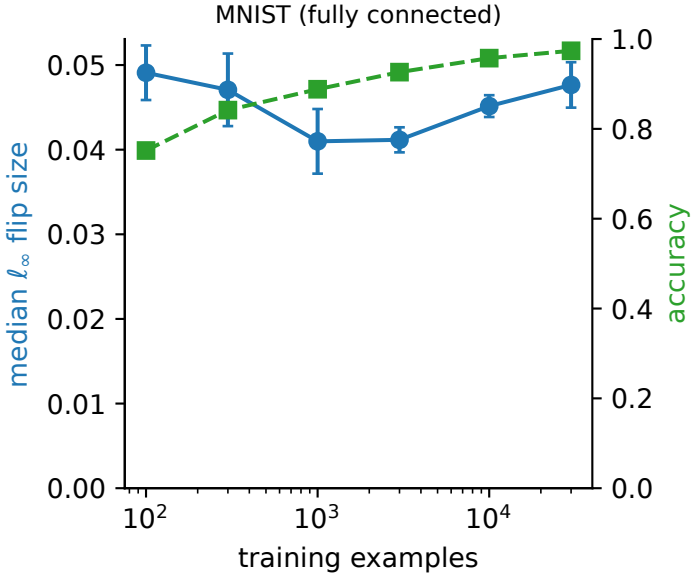
MNIST (fully connected)
median ℓ_∞ flip size
accuracy
training examples
0.00
0.01
0.02
0.03
0.04
0.05
0.0
0.2
0.4
0.6
0.8
1.0
10^2
10^3
10^4

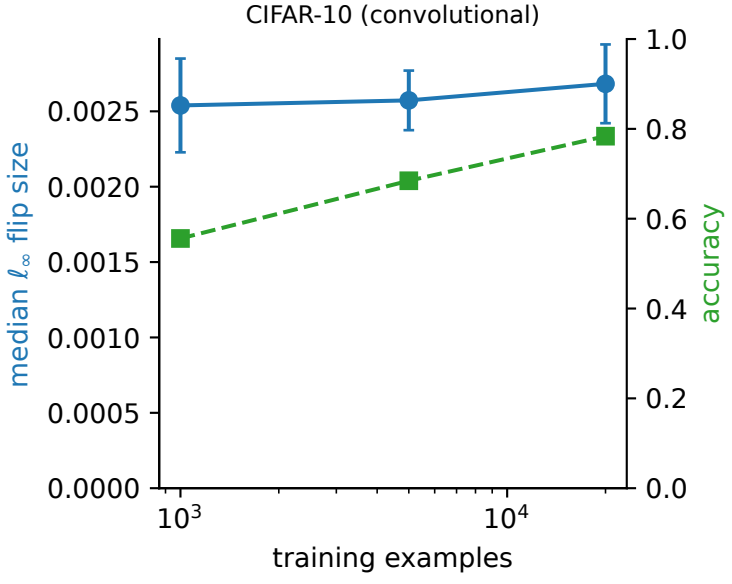
CIFAR-10 (convolutional)
median ℓ_∞ flip size
accuracy
training examples
0.0000
0.0005
0.0010
0.0015
0.0020
0.0025
0.0
0.2
0.4
0.6
0.8
1.0
10^3
10^4

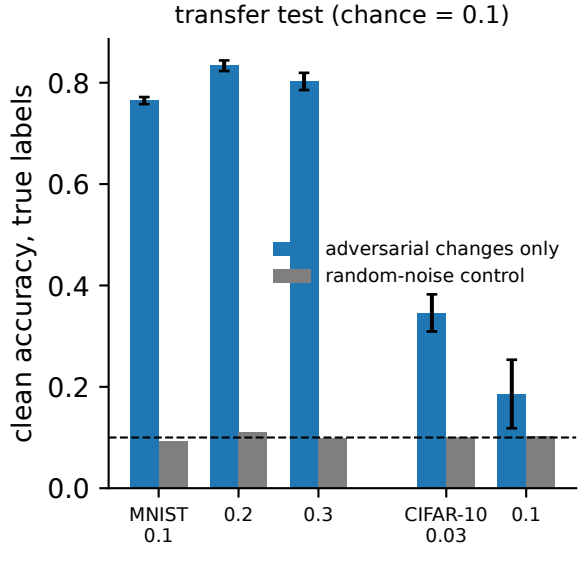
transfer test (chance = 0.1)
clean accuracy, true labels
adversarial changes only
random-noise control
0.0
0.2
0.4
0.6
0.8
MNIST
0.1
0.2
0.3
CIFAR-10
0.03
0.1